The MOOS-IvP Build System

Christian Convey (christian.convey@navy.mil)

2007-08-14

Abstract

This document details the design and intended use of MOOS-IvP's build system.

Contents

1   HOWTO Build MOOS-IvP

This section gives some basic background in using any CMake-based build system, and then explains how to use MOOS-IvP's particular build system.

1.1   CMake

1.1.1   Overview

This section gives a brief overview of CMake, but clearer and more complete sources of information exist:

- The project's website: `cmake.org`

- <u>Mastering CMake</u>, by Ken Martin and Bill Hoffman.

- The CMake users email list: `cmake@cmake.org`. You can sign up for this at the `cmake.org` website.

CMake (`cmake.org`) is a cross-platform build system. Compared to traditional Makefiles, CMake's control files (named `CMakeLists.txt`) tend to be much shorter and easier to read.

CMake has a number of back-ends for various operating systems and Make programs (gmake, nmake, etc.) and IDEs (MS VisualStudio, KDevelop, Apple's XCode, etc.) CMake-based build systems trivially support some features that can be a real hassle to implement in hand-written Makefiles, such as calculating header file dependencies and supporting an `install` target.

The `cmake` program reads in a project's `CMakeLists.txt` files and produces one or more files suitable for use in your build-system of choice (GNU Make, MS VisualStudio, etc.) These generated build files will typically invoke the `cmake` program when building certain targets; therefore CMake must be installed on any computer that will execute `cmake`-produced build files.

1.1.2 Language for Describing Makefiles

CMakeLists.txt files are written in an imperative language which is documented on CMake's website. The language lets you set variables and supports branching, looping, and a limited form of subroutines (called MACRO's).

When cmake processes a project's CMakeLists.txt files, the program embodied in those files specifies the details of the build system to be created. The CMakeLists.txt files execute sequentially, specifying details as they run. Then end product is not a fully built version of your project, but rather a set of build files (Makefile, etc.) that when executed will build your project.

The sequential nature of the programming language lets variables be computed, helper programs run, conditions to be tested, etc. in order to determine what should happen when the person building the project finally runs make (or performs the equivalent action in VisualStudio, XCode, etc.).

1.1.3 Federated Builds

A typical CMake-based build system will contain one CMakeLists.txt file in each of the project's source tree's directories. A simplified depiction of this scheme is shown in Figure 1.

The project's *top-level* CMakeLists.txt file typically sets variables that govern the entire project's build details, such as the directory into which the built libraries and / or executables should be placed after being linked, and the directory(ies) in which the project's header files can be found.

The CMakeLists.txt files in the *internal levels* of the source tree often do nothing other than call ADD_SUBDIRECTORY on each of that directory's subdirectories. (ADD_SUBDIRECTORY is very similar in effect to the C preprocessor's #include directive.) The practice of having intermediate-level directories' CMakeLists.txt files mostly just consist of ADD_SUBDIRECTORY is common because those directories often exist merely to group together subdirectories into units that are meaningful to the programmer. Therefore those directories have little significance to the way the project code is actually built.

The source tree's *leaf directories* typically contain the bulk of the project's interesting source code, with each subdirectory having the code for a single library or executable. The CMakeLists.txt file in one of these leaf directories will generally contain an ADD_LIBRARY or ADD_EXECUTABLE command which leads to the compilation and linking of that one particular program / library.

Figure 1: Conventionally each directory in the source code tree contains a CMakeLists.txt file.

1.1.4 Inherited Variables

`CMakeLists.txt` scripts have variables, similar in form and purpose to `make` shell variables. They're treated and manipulated as text, and lists are generally represented as a single string in which spaces separate the list's elements. (Another convention is sometimes used in CMake's functions, where semicolons rather than spaces separate list elements.)

In traditional `Makefile`s, a child `make` process inherits a parent `make` process' variable values. A similar effect exists in a hierarchy of `CMakeLists.txt` files, but via a different mechanism:

Consider Figure 2. Using these `CMakeLists.txt` files would look like this:

```
cmake -f CMakeLists.txt
(... some miscellaneous output from cmake ...)
Y
X
(... some miscellaneous output from cmake ...)
```

Notice how the subdirectory's file `Bar/CMakeLists.txt` not only inherits a value of `MyVar` from the parent `CMakeLists.txt` file, but `Bar/CMakeLists.txt` also affects the behavior of the parent `CMakeLists.txt` file. (`Bar/CMakeLists.txt`) modifies the value of `MyVar`, and that new value is in effect in the lines of the parent `CMakeLists.txt` file that follow the `ADD_SUBDIRECTORY` call.

Contrast this to behavior to a `Makefile`-based system where a parent directory's `Makefile` calls `make -f Bar/Makefile`. In that scenario, there is no mechanism by which `Bar/Makefile`'s modifications of `make` variables could propagate back to the parent directory's running instance of `make`.

1.1.5 Cache Files

Some CMake variables are computed each time you run `cmake`, but other *cached* variables, once set, have their values recorded and stored in a file named `CMakeCache.txt`. Variables stored in this file can be edited using the `ccmake` program (Linux) or the `CMakeSetup.exe` program (Windows). This not only lets users enter certain data just one time (i.e., the string that specifies the version of MOOS being built), but it also exposes some CMake variables that whose value the user might want to correct.

For example, when CMake looks for the FLTK libraries, one variable that gets set is `FLTK_DIR`. If CMake's `FindFLTK` package can't figure out where the FLTK libraries are, it will set the cached `FLTK_DIR` variable to have a value of `FLTK_DIR-NOTFOUND`. The user can run `ccmake` or `CMakeSetup.exe` to manually set the value of `FLTK_DIR` to a valid value (typically `/usr/lib` for that particular

Figure 2: ADD_SUBDIRECTORY joins CMakeLists.txt files much like C's include

Figure 3: Build-time relationship between CMake, make, and their control files

variable), and then re-run `cmake`. This is typically enough of a help for a CMake package such as `FindFLTK` package to work out the rest of the details it needs to.

Sometimes a `CMakeCache.txt` file will contain so many bad values that the user decides it's better just to start from scratch. In general it's safe for a user to manually delete a `CMakeCache.txt` file. When he next runs `cmake`, the variables whose values had been contained in the file will simply be freshly computed.

1.1.6 cmake, ccmake, CMakeSetup.exe

Figure 3 shows the relationship between CMake, the traditional `make` program, and various control files. `cmake` is used to generate build-system files from a set of `CMakeLists.txt` files. `cmake` will also read from and write to a `CMakeCache.txt` in the directory from which `cmake` is run.

Sometimes `cmake` doesn't succeed the first time it's run. The FLTK module from above is a good example. We can easily write `CMakeLists.txt` that would

terminate the `cmake` invocation if the cached `FLTK_DIR` variable had a value of `FLTK_DIR-NOTFOUND`.

One solution to such a problem would be for the user to specify a good value on the `cmake` command line. For example:

```
> cmake -DFLTK_DIR=/usr/lib -f CMakeLists.txt
```

This approach is sometimes the best, but `ccmake` and `CMakeSetup.exe` offer a more interactive approach to configuring the build system. Within these tools, you continue re-running `cmake` and modifing cached variable values until the `cmake` succeeds. At that point the user instructs the tool to "generate" the build system files (on Linux, this typically includes a set of `Makefile`s). Once this is done the user will be ready to run `make` (or whatever the equivalent action is for his build tool).

Note that `ccmake` and `CMakeSetup.exe` will not create the build system files until the set of CMake variable values has stabilized. This means that the tool must observe that when it runs `cmake`, none of the cached variables' values have changed. This relates to the fact that run n of `cmake` may set new value in a cached variable, but the ramifications of that new variable aren't fully realized until the next invocation of `cmake` makes use of that value.

Practically speaking the need to iteratively run `cmake` within `ccmake` and `CMakeSetup.exe` doesn't feel burdensome. And if one knows the proper values for all of the CMake variables of interest, you can set them on the `cmake` command line as shown above, thus entirely avoiding the need to run `ccmake` or `CMakeSetup.exe`.

1.1.7   Finding External Packages

When it comes to telling the build system where to find the header files and development libraries for external packages such as FLTK or libGLU, CMake is similar to GNU's `autoconf`: Contributors write CMake modules to discover the details needed to use a particular program, library, etc. `automake` packages are written in the `m4` macro language, whereas CMake modules are written in CMake's scripting language - the same language used in `CMakeLists.txt` files.

CMake includes a number of these packages as a part of its standard distribution. On Ubuntu Linux you can see those modules in the directory

```
 /usr/share/cmake-2.4/Modules/
```

You can also write your own modules, but you need to tell CMake about which directories to search when looking for extra modules that you've written.

Sometimes the standard supplied modules are of lower quality and you might want to use your own CMake code instead. A good example is `FindFLTK`, which is used by both MOOS and IvP in order to discover the locations of the include files and libraries for the FLTK library. `FindFLTK` does a bad job of discovering the locations of these files, even if they're in `/usr/include/FL` and `/usr/lib`, as expected. `FindFLTK` is also structured so that it's inconvenient to test whether or not FLTKs header files and libraries have been found, without also requireing that the `fluid` program (unused by MOOS and IvP) is installed.

MOOS avoids this quality problem by using a customized version of the module, implemented in the file:

```
MOOS/MOOSFindFLTK.cmake
```

But even that customized version of `FindFLTK` has a painfully complex implementation, and it still needs help just to realize that it should look in the directory `/usr/lib` to find the FLTK libraries.

1.1.8   Multiple Attempts to Configure the Project

Configuring a CMake-based project is the act of attempting to create the build files (i.e., `Makefile`, etc.) needed to build your software. As described in Section 1.1.6, this make require multiple attempts in order to ensure that all of the important CMake variables have the values required to build your software the way you want it.

1.1.9   In-source vs. Out-of-source Builds

Sometimes we want to maintain multiple builds of the same software, where the builds differ by details such as debug vs. optimized build, or perhaps where files will be copied to if one were to run the command `make install`. (In the MOOS-IvP project, the scripts for creating the Debian package files rely on this feature.)

To support the concurrent existence of multiple kinds of builds, the binaries produced by the build process (object files, libraries, and programs) for different types of builds must reside in different directories. Even the generated Makefiles for the different build types must be separated, because (for example) the directories to which files will be copied when `make install` is run are stored in the Makefiles generated by CMake.

An *in-source build* is one where the Makefiles, object files, etc. associated with a build of the software are stored in the same directory tree as the source files themselves. An *out-of-source build* is one where the Makefiles generated by CMake, and the object files, libraries, and programs produced by those Makefiles, are stored in a different directory tree than that in which the source

code resides.

CMake can accomodate either in-source or out-of-source builds. When one creates an out-of-source build, CMake will generally create a directory tree whose structure mirrors that of the source tree, and populate that newly created directory tree with Makefiles and various extra subdirectories needed to fully build the software. When one creates an in-source build, CMake will place those same Makefiles and extra subdirectories directly into the appropriate locations within the source code's directory tree.

When one runs `cmake` or `ccmake` or `CMakeSetup.exe` to initially configure a project, he passes as a command-line argument the root of the source code directory tree. CMake will look within that specified directory tree for the `CMakeLists.txt` files that guide the creation of the build files, but it will create the build files (`Makefile`, etc.) in the current working directory from which `cmake` was invoked. Therefore an in-source build tends to look like this:

```
cd ~/MyProj/src
ccmake ./
make
```

and an out-of-source build tends to look liks this:

```
mkdir ~/MyBuildDir
cd ~/MyBuildDir
ccmake ~/MyProj/src
make
```

It's often preferable to use out-of-source builds when possible. When one configures a project using `cmake` (or `ccmake` or `CMakeSetup.exe`), certain details become fixed and cannot be changed no matter what environment variables you set prior to running `make` within that build tree. Examples include whether the code will be built in debug or release mode (governed by the value of the `CMAKE_BUILD_TYPE` variable), and where files will be installed if you run `make install` (governed by the value of the `CMAKE_INSTALL_PREFIX` variable).

If one wants to rapidly switch, for example, between release and debug builds, then it may be desirable to maintain two separate out-of-source configurations that differ only in their value of the `CMAKE_BUILD_TYPE` variable. One might therefore do something like this:

```
mkdir ~/MyBuildDir/release
cd ~/MyBuildDir/release
ccmake -DCMAKE_BUILD_TYPE=Release ~/MyProj/src
make
```

```
mkdir ~/MyBuildDir/debug
cd ~/MyBuildDir/debug
ccmake -DCMAKE_BUILD_TYPE=Debug ~/MyProj/src
make
```

Another example of where the use of out-of-source configurations is useful is when the MOOS-IvP build system produces Debian packages. Our technique for producing the different Debian packages involves varying the value of the `CMAKE_INSTALL_COMPONENT` CMake variable. It seemed overly intrusive to re-configure the build directory used by the person creating the Debian packages, so the Debian packaging scripts use out-of-source builds whose configurations have various values for the `CMAKE_INSTALL_COMPONENT` CMake variable.

In order to remain consistent with the general behavior of the of IvP's pre-vious build system, the top level `build.sh` script performs an in-source build. However, the MOOS-IvP project can be built equally well as an out-of-source build.

1.2   Source Tree Organization and the Two Build Systems

When targeting a Make-based build system, `cmake` will produce a set of files named `Makefile`, one in each directory of the tree.

In order to permit MOOS-IvP's old set of Makefiles to remain present in the source tree, I renamed all of the old Makefiles to be named `makefile` (note the lower-case filename). In the old make system, the makefile in one directory builds the contents of its child directories by invoking the command `make -f (child-dir)/makefile`. Because the tweaked version of the old build system specifically calls `make` on the lower-case-named `makefile`, the two build systems manage to remain completely independent.

Our expectation is that after several months of gaining confidence in the CMake-based build system, Mike Benjamin will permit the old `makefile` files to be removed from the source tree.

1.3   Building MOOS-IvP

1.3.1   Required Steps for old makefile-based build

In contrast to a previous organization of the MOOS-IvP directory tree, the `ivp/src` no longer contains a subset of the programs and libraries that are part of the MOOS project. Therefore if one is using the old makefiles to build IvP, he must first explicitly build the MOOS libraries so that they're available when the ivp programs are being linked. This is the primary difference between the old and new technique for building IvP when using the old-style makefiles.

The steps are as follows:

1. Check-out or download the moo-ivp source code, into some directory `foo/`.

2. Build MOOS:

   2.1. `cd foo/MOOS`

   2.2. `../build.sh`

   If this step fails, you may need to modify one or more of the directories mentioned in the file `../build.sh`.

   Note that MOOS is being built with CMake, even though these present instructions describe building IvP without using CMake. This is because (a) we've decided to stop importing code from the MOOS project directly into the IvP source tree, and (b) the MOOS project uses CMake rather than hand-written makefiles.

3. `cd foo/ivp/src`

4. `make -f makefile`

   If you don't specify `-f makefile`, and have previously used the CMake-based build system in this directory, then your invocation of `make` will probable use the file `ivp/src/Makefile` (note the upper-case `M`). That file is created by CMake, which presumably isn't what you want since these present steps describe how to use IvP's *non*-CMake build system.

1.3.2 Required Steps for CMake-based build (Easy Technique)

1. Check-out or download the moo-ivp source code, into some directory `foo/`.

2. `cd foo/`

3. `./build.sh`

1.3.3 Required Steps for CMake-based build (Advanced Technique)

This is the approach you'll need to take if CMake can't find some required header file or library, or if you want to tweak what set of programs and libraries are built.

1. Check-out or download the moo-ivp source code, into some directory `foo/`.

2. `cd foo/`

3. `ccmake -f CMakeLists.txt`

   Within `ccmake`, set variables as desired until you can generate the build files. Then generate them (within `ccmake`). Documentation of `ccmake` and `CMakeSetup.exe` is out of the scope of this document, but the book Mastering CMake is an excellent source of information. An explanation can also be found here:

   `http://www.cmake.org/HTML/RunningCMake.html`

   Note that the file `foo/build.sh` shows an invocation of `cmake` that sets some CMake variables that the current build system has trouble figuring out on your own. If you're having trouble getting any of those variables to take on good value when using `ccmake`, consider using the respective values mentioned in `foo/build.sh`

4. `make`

1.3.4   Requried Ubuntu Packages

The following is a (perhaps) partial list of packages that must be installed on a Ubuntu 7.04 ("Feisty Fawn") system in order to build IvP and the non-Matlab parts of MOOS:

Ubuntu package	Reason needed
cmake	Integral part of build system
mesa-common-dev	Requried to build programs that use FLTK + OpenGL
libfltk1.1-dev	Requried to build programs that use FLTK + OpenGL
libglu1-mesa-dev	Requried to build programs that use FLTK + OpenGL
libtiff4-dev	Requried to build programs that use FLTK + OpenGL
libxft	Requried to build programs that use FLTK + OpenGL
libxft-dev	Requried to build programs that use FLTK + OpenGL
fluid	To satisfy CMake's `FindFLTK` module
ncurses-dev	Required to build MOOSRemote
python-dev	Required to build `uVoice` (not built by default)
swig	Required to build `uVoice` (not built by default)

1.4   Important CMake Variables

1.4.1   For FLTK

FLTK_DIR - This is the library where CMake's `FindFLTK` module expects to find the FLTK libraries. For mystifying reasons, it can't figure out on its own to look in `/usr/lib`. You will probably want to set this to `/usr/lib`.

As a bonus, once you do so, MOOS's customized version of that module, `MOOSFindFLTK`, will be able to figure out the proper value for `FLTK_INCLUDE_DIR` all by itself. As of this writing, IvP uses the standard version of `FindFLTK` which sadly cannot figure out the proper value for `FLTK_INCLUDE_DIR` after you've provided a valid value for `FLTK_DIR`.

FLTK_INCLUDE_DIR - The directory that is expected to include a subdirectory named `FL`, which in turn contains FLTK's source files. On Ubuntu Linux 7.04, at least, the proper value for this is `/usr/include`.

1.4.2   For Python.h

IVP_BUILD_MOOS - When this is `ON` the build system will build `uVoice`, a Python wrapper for MOOS's libraries. When this is `ON` you'll need to set a proper value for the CMake variable `PYTHON_INCLUDE_PATH` or else the build processes won't attempt to build `uVoice`.

PYTHON_INCLUDE_PATH - The directory containing the version of `Python.h` that should be used when building `uVoice`. On Ubuntu Linux 7.04, which by default uses Python version 2.5, the correct value for this variable is `/usr/include/python2.5`, assuming that you're trying to make a wrapper that works with the host system's version of Python.

1.4.3   Other CMake Variables

IVP_BUILD_... - There's one of these variables for each program and library in the IvP part of the moos-ivp source tree. Setting a value to `ON` will cause the corresponding program or library to be included in the build process.

Several programs / libraries that were by default *not* built by the old build system have a default value of `OFF` for this variable. All others default to `ON`.

Note that the current IvP build system does not enforce, at this step, the build dependencies between libraries and the programs that require them at link-time. If you disable the building of a library but enable the building of some program that must be linked to that library, you'll end up with an error from the linker when that program is being built.

BUILD_... - These come from the MOOS side of the project. Setting one of these variables to `ON` may lead to the appearance of even more such variables when you next configure the build tree, because MOOS's build system supports a two-level hierarchy of dis-/en-abling the building of its programs and libraries.

CMAKE_BUILD_TYPE - By convention, any CMake variable whose name starts with `CMAKE_` has special meaning to CMake. This variable controls details related to debug vs. optimized builds. Two noteworthy values are `Debug` and `Release`.

2 MOOS-IvP Build System Implementation Details

2.1 create-subdir-cmake-files.sh

IvP's source tree contains about 79 libraries and applications as of this writing, each of which requires its own `CMakeLists.txt` file. Most IvP programs have similar needs for their `CMakeLists.txt` files, and the same goes for most IvP libraries.

Because this regularity exists, and because all 79 of those `CMakeLists.txt` files must be rewritten when certain improvements are made to the build system, IvP contains a script (plus a few helper files) to re-create those 79 `CMakeLists.txt` files whenever required.

That script is named `create-subdir-cmake-files.sh`. It and its support files are found in the directory

`ivp/scripts/util/create-subdir-cmake-files/`

Whenever a library or program is added to or removed from IvP, this script should be modified so that we can always easily recreate and / or redesign IvP's build system.

As of this writing, this script only creates IvP's `CMakeLists.txt` files, and none of MOOS's. This is because the MOOS project already has a CMake-based build system in place and there's no obvious need to re-implement it.

3 Future Work

3.1 Debian packaging

Debian packaging can be accomplished by some minor additions to IvP's CMake-based build system. We've already submitted a patch for doing this in MOOS to Paul Newman, so the techniques and code have been worked out already.

3.2 RPM packaging

CMake 2.6, expected to be released in late 2007, should have a feature for easily producing RedHat Package Manager (RPM) software packages. When this happens we may support this for MOOS-IvP.

3.3 Windows builds

Now that all of MOOS-IvP can be built with CMake, we may be able to build the software on Windows with little or no additional effort.

[Note: Andrew Shafer has protested that this may be an overly optimistic assesment.]

3.4 Windows packaging

Develop an installation technique for MOOS-IvP that's convenient for Windows programmers. We haven't yet investigated what the most preferred form of packaging is.

3.5 Building Documentation

CMake can easily build man pages, Doxygen documentation, etc. Building it with CMake helps us to package the documentation for end-user use, as well as creating `...-doc` Debian packages and whatever the RPM equivalent would be.

3.6 Building IvP's Website

CMake's scripting language should be sufficient for expressing how to create new versions of IvP's static html website. This would be useful to ensure that the website is always synchronized with the most current release of our software and documentation.